

Search of Deep Learning Models based on Asymmetric Adaptivity and Locality Sensitive Hashing

Lixi Zhou
Arizona State University
Tempe, USA
lixizhou@asu.edu

Zijie Wang
Arizona State University
Tempe, USA
zijiewang@asu.edu

Jia Zou
Arizona State University
Tempe, USA
jia.zou@asu.edu

ABSTRACT

In recent, deep learning has become the most popular direction in machine learning and artificial intelligence. However, preparation of training data, as well as model training, are often time-consuming and become the bottleneck of the end-to-end machine learning life-cycle. Reusing models for inferring a dataset can avoid the problem and greatly reduce the model deployment costs. However, it is challenging to discover the right model for reuse. Although there exist a number of **model sharing platforms such as ModelDB, TensorFlow Hub, PyTorch Hub, DLHub, most of these systems require model uploaders to manually specify the details of each model and model downloaders to screen keyword search results for selecting a model. We are lacking a highly productive model search tool.** This paper proposes an end-to-end model search process by matching the target query dataset with the training datasets of the available models. We identified a significant challenge that the matching relationship between an target query dataset and a training dataset is asymmetric. Therefore, we proposed a new adaptivity metric, and a novel two-level locality sensitive hashing index to facilitate the comparison of the adaptivity of the target query dataset to a number of models' training datasets in order to search for the proper model to reuse. We demonstrated the effectiveness of the proposed approach using four different applications: activity recognition, entity matching, image recognition, and natural language processing.

ACM Reference Format:

Lixi Zhou, Zijie Wang, and Jia Zou. 2018. Search of Deep Learning Models based on Asymmetric Adaptivity and Locality Sensitive Hashing. In *Woodstock '18: ACM Symposium on Neural Gaze Detection, June 03–05, 2018, Woodstock, NY*. ACM, New York, NY, USA, 9 pages. <https://doi.org/10.1145/1122445.1122456>

1 INTRODUCTION

Pre-trained models can be reused directly or through transfer learning so that the substantial efforts required for collecting and annotating training data on the new problems can be greatly saved [38]. The prediction of the serving accuracy of a model in a target query dataset is inherently a complex problem. It is related to a number

of factors, such as the model architecture, initial weights, hyper-parameters, and most importantly, the similarity and adaptivity between the source training dataset and target query datasets [18]. In this work, we focus on a broad class of model versioning scenarios, where the users need to choose a few models from a large repository of models that have similar architectures.

Motivating Scenario 1. Marketplace-level Forecasting at Uber.

They forecast supply, demand, and other quantities in real time for hundreds of cities across the globe. Uber is operating in many cities across the world, and different cities may pose different geospatial characteristics [34]. Besides, the Uber business might be at different growth stages for different cities. Therefore, they shard the problem spatially by city, training a model instance for each city-quantity combination using the same model architecture. In addition, they frequently update the models through retraining, when they detect model performance degradation due to the changing market conditions, and they need to independently trigger the retraining of the models for a city. All of these leads to thousands of or even more model versions [34].

Motivating Scenario 2. Downtime Prediction at Manufacturers.

A steel plate manufacturer [37] owns several plants in different countries. The data of machines from different locations are collected centrally. This includes time-series data, such as temperature, humidity, and vibration data, as well as high-resolution image data from cameras. Bob wants to create a machine learning model that predicts the downtimes for a specific machine type. However, the failure situation is changing from time to time, therefore Bob retrains and redeploys the model from time to time to keep it up-to-date, thus leading to many different versions.

Model versioning has motivated a number of model management platforms, such as Gallery [34], ModelDB [35], ModelHub [28], Metadata Tracking [32], MLFlow [12], TensorFlow Hub [4], PyTorch Hub [3], and DLHub [1]. However, the discovery of the proper model versions for serving is mostly through the intuition of domain experts and trial-error processes. Most of these platforms provide a model search tool based on the model metadata as shown in Tab. 1.

Table 1: Model search based on model metadata

Search Category	Filter Examples
Publication Schemas	author, date, description
Model Information	algorithm, software version, network architecture, dependency
Development Provenance	versions, contributors
Training Information	training datasets, parameters
Performance	accuracy, recall, precision, evaluation dataset, evaluation date

However, none of the above information can directly tell *which version(s) of a model should be selected for inference on a new dataset*, which is exactly the problem that we will address in this study. We observed that existing model serving platforms lack an efficient

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

Woodstock '18, June 03–05, 2018, Woodstock, NY

© 2018 Association for Computing Machinery.

ACM ISBN 978-1-4503-XXXX-X/18/06...\$15.00

<https://doi.org/10.1145/1122445.1122456>

mechanism to search for a few model versions that will have the best inference accuracy on the target query dataset. Instead, users need to perform a text-based search using keywords, scrutinize each search result by deploying and testing the model. Such human-centered model selection based on trial and error is inefficient, which delays the model deployment and incurs significant human costs. To alleviate such overheads and human costs, it is urgent to automate this process.

A basic assumption in learning theory is that the training and testing sets are drawn from the same probability distribution [18]. Therefore, if the pre-trained model's source domain (i.e., the domain of the training data) has similar probability distribution with the target domain (i.e., the domain of the testing data), the pre-trained model may achieve good accuracy in the target domain. This motivates a nearest neighbor search approach that selects the models of which the training dataset are most similar to the target query dataset. However, there exist several **challenges** for this approach:

First, the inference accuracy is asymmetric. For example, a dataset A may contain only a subset of instances of another dataset B; then supposing we have a model M_A that is trained on A, and a model M_B that is trained on B, the accuracy of M_A on B, and the accuracy of M_B on A could be very different. Because of such asymmetry, it is problematic to simply apply the symmetric similarity measurement to the features of the training dataset and the features of the target query dataset.

Second, given a large number of models and large size of training data (as illustrated in Fig. 1), pair-wise comparison of all models' training datasets to the target query dataset is not practical for high-frequency model search requests [34]. In addition, search response times increase from 1 to 4 seconds, user experience as well as conversion rates drop sharply [2, 5].




Figure 1: TensorFlow Hub Models. The value and size of the circle denote the number of models trained on the dataset.

Main Ideas To address the problem regarding the asymmetric serving accuracy on the swapping of source and target domains, we propose a novel similarity measurement called **adaptivity**. It first requires to partition the training dataset and the target query dataset into multiple partitions of same size respectively. Then, the adaptivity is computed as the set containment of the training dataset to the target query dataset, which is defined as the ratio of the number of the similar partitions to the number of the partitions in the target query dataset.

To reduce the pair-wise comparison overhead, we consider leveraging Locality Sensitive Hash (LSH), which is a popular technique

to solve nearest neighbor problems. Most of the existing LSH families, are designed to measure the Hamming distance [15], Euclidean distance [19], and cosine similarity [10] of two vectors with fixed dimensions, or the Jaccard similarity [8] for two high-dimensional binary vectors or two sets of items. Until very recently, the LSH for measuring similarity of probability distributions is studied [13, 27]. However, these LSH schemes are not designed for the adaptivity metric that we propose.

A great benefit of the adaptivity metric is that it can be converted to Jaccard similarity [8], which motivates us to use the Minwise LSH to accelerate the pair-wise set containment computations. However, the problem is, Minwise LSH is designed for indexing high-dimensional vectors, but in our case, each dataset is a vector of partitions and it is hard to directly derive the Minwise LSH for such a complicated vector, lacking an efficient comparator for two partitions. To address the problem, we further propose a novel two-level LSH index. The top level is the aforementioned Minwise LSH index, and the bottom level is a LSH index for quickly identifying similar partitions that have small JS-divergence with each other, which we call as JSD-LSH.

Utilizing the JSD-LSH, each partition is converted to a JSD-LSH signature, and two partitions have the same signature if they have similar probability distributions (i.e., small JS-divergence). By seeing each signature band as a value, a dataset is abstracted as a vector of values, from which Minwise LSH can be easily computed for estimating the adaptivity.

Our Key Contributions are summarized as follows:

- (1) In this work, we identified the problem of searching related models for serving a target dataset from a set of models that are trained using similar methodology with different training data. Our work will greatly enhance the capability of existing model repositories and improve the productivity of the model reuse process, which is urgently needed in production.
- (2) We proposed a novel asymmetric adaptivity measurement for a pair of a training dataset and a target query dataset to evaluate how well the model that is trained using the training dataset will work on the target query dataset.
- (3) We proposed a unique two-level LSH index leveraging LSH mechanisms for JS-divergence and Jaccard similarity at different levels to improve the similarity comparison efficiency with theoretical proofs.
- (4) We implemented the proposed approach and conducted extensive experiments to compare the effectiveness and efficiency of the proposed approach to alternative similarity measurements and model search strategies using four different applications: activity recognition, entity matching, image recognition, and natural language processing.

2 BACKGROUND

2.1 Jensen-Shannon (JS) Divergence

JS-divergence [25] is a measurement of the similarity between two probability distributions, which is a symmetric metric derived from the asymmetric Kullback-Leibler divergence [21].

Let P and Q be two probability distributions associated with a common sample space Ω , and let $M = (P + Q)/2$. The JS-divergence is defined by:

$$D_{JS}(P \parallel Q) = \frac{1}{2}D_{KL}(P \parallel M) + \frac{1}{2}D_{KL}(Q \parallel M) \quad (1)$$

Here, D_{KL} denotes the Kullback-Leibler divergence, which is defined as following for discrete distributions:

$$D_{KL}(P \parallel Q) = \sum_{\mathbf{x} \in \Omega} P(\mathbf{x}) \log\left(\frac{P(\mathbf{x})}{Q(\mathbf{x})}\right), \quad (2)$$

and as following for continuous distributions:

$$D_{KL}(P \parallel Q) = \int_{\Omega} P(\mathbf{x}) \log\left(\frac{P(\mathbf{x})}{Q(\mathbf{x})}\right) d\mathbf{x} \quad (3)$$

In this work, we mainly consider leveraging JS-divergence to tell the similarity of two partitions. The reasons are: First, JS-divergence is a widely used similarity metric for probability distributions [22]. Second, it is easier to find LSH schemes for JS-divergence, compared to other similarity measurements of probability distributions [27].

While other metrics such as Maximum Mean Discrepancy (MMD) could measure the domain adaptivity more accurately, it requires an optimization process and significantly higher computational costs, as illustrated in Tab. 2.

Table 2: Average latency comparison on Activity Recognition (Sec. 4.1) datasets (Unit: seconds).

Jaccard similarity	KL-divergence	JS-divergence (w/o LSH)	MMD
379	29	87	6451

2.2 Locality Sensitive Hashing (LSH)

LSH was developed for general approximate nearest neighbor search problem [19]. It requires a family of LSH functions, each of which is a hash function whose collision probability increases with the similarity of the inputs. The (r_1, r_2, p_1, p_2) -sensitive LSH family is formally defined in Definition. 2.1.

Definition 2.1. Let $\mathcal{F} = \{h : M \rightarrow U\}$ be a family of hash functions for distance measurement D . \mathcal{F} is (r_1, r_2, p_1, p_2) -sensitive ($r_1 < r_2$ and $p_1 > p_2$), if $\forall p, q \in M$, it satisfies that: (1) if $D(p, q) \leq r_1$, we have $Pr[h(p) = h(q)] \geq p_1$; (2) if $D(p, q) \geq r_2$, we have $Pr[h(p) = h(q)] \leq p_2$.

LSH is first proposed by Indyk and et al. [19] for measuring the Hamming distance in a d -dimensional Euclidean space, which requires the data to be vectors with fixed dimensions. MinHash [8] is a widely used family of hash functions for Jaccard similarity. SimHash is an LSH scheme for Cosine distance [10]. Both Minwise LSH and SimHash are only applicable to high-dimensional binary vectors or sets of values (without fixed dimensions), and MinHash is usually considered to be more computationally efficient than SimHash [43].

2.2.1 Minwise LSH. Minwise LSH [8] was proposed as an index for estimating the Jaccard similarity between two sets of values. It converts each set into a MinHash signature using a family of Minwise hash functions. Each Minwise function hashes each value in the set into an integer, and returns the minimum hash value.

Given a Minwise hash function h_{min} , it has been proved that the probability of the two minimum hash values, which are observed over two sets of values X and Y , equal to each other, is the Jaccard similarity of X and Y : $P[h_{min}(X) = h_{min}(Y)] = jaccard(X, Y) = \frac{X \cap Y}{X \cup Y}$ [8]. Furthermore, usually a family of K Minwise hash functions are applied to a set of values, generating K minimum hash values as a signature. Given the signatures of X and Y , Jaccard similarity can be estimated as the ratio of the number of collisions in the corresponding minimum hash values to K . In practice, for convenience, the K values are further partitioned into L bands, so that the Jaccard similarity can be estimated as the ratio of the number of bands from the two sets that are equivalent to the total number of bands. The lookup of matching bands can be accelerated by storing the Minwise signatures of all candidates corresponding to each band into a hash table.

2.2.2 LSH for JS-divergence. As mentioned, the LSH schemes for probability distributions are recently studied [13, 27]. S2JSD-LSH [27] provides LSH functions for a measurement that approximates the square root of two times the JS-divergence. Unfortunately, their LSH design doesn't provide any bound on the actual JS-divergence. Then Chen et al. [13] propose a new LSH scheme for the generalized JS-divergence through the approximation of the squared Hellinger distance, which is proved to be bounded with the actual JS-divergence and thus used in this work, as defined in Eq. 4:

$$h_{a,b} = \lceil \frac{\mathbf{a} \cdot \sqrt{P} + b}{r} \rceil, \quad (4)$$

, where P is a probability distribution in the sample space Ω , $\mathbf{a} \sim \mathcal{N}(0, I)$ is a $|\Omega|$ -dimensional standard normal random vector, \cdot represents the inner product operation, b is a random variable uniformly distributed on $[0, r]$, and r is a positive real number. This approximation is proved to be lower bounded by a factor 0.69 for the JS-divergence [13].

3 METHODOLOGY

In this section, we first explain why symmetric metrics that measure the similarity of two sets of features cannot be used to estimate whether a model trained in one set will effectively infer the other dataset. Then, to address the issue, we propose a new asymmetric adaptivity. We further propose a novel two-level LSH index framework to accelerate the system-wide computation of adaptivity with theoretical proofs.

3.1 Problem Definition

We are given a list of models of similar architecture, each of which is trained using a different training dataset that shares the same schema of features and labels (e.g., generated from similar data pipelines in the same organization). The set of training datasets for these models is denoted as $T = \{t_1, \dots, t_k\}$. A training dataset for the i -th model is denoted as $t_i = \{\langle \vec{x}, y \rangle\}$, where each training sample consists of a feature vector \vec{x} , and a label y . Similarly, a target dataset is denoted as $\mathbf{I} = \{\langle \vec{x}', y' \rangle\}$, where each testing sample consists of a feature vector \vec{x}' , which is known, and a label y' , which is unknown and needs to be inferred. The schema of the target dataset is consistent with the training datasets. The question

is which model should be selected so that it will achieve the best inference accuracy on I .

Why symmetric similarity metrics will not work for this problem? A straightforward idea is to simply measure the similarity between the set of features to be inferred, denoted as $F_I = \{\{\vec{x}'\}\}$, and the set of features of each training dataset, denoted as $F_i = \{\{\vec{x}\}\}$, and select the model that has the most similar training dataset. However, if the similarity measurement is symmetric, there exists a significant problem that the results will be skewed by the discrepant sizes of the training dataset and target dataset. For example, if $F_I \subset F_i$, the model trained on t_i may have good accuracy on I , because the model has seen all of the samples in I during the training process. On the contrary, the model trained on I may have relatively worse accuracy on F_i , because the model has only seen a portion of samples in F_i in training. Therefore, it's more important to measure the ratio of the similar samples to the samples in the target dataset, which is asymmetric.

3.2 Adaptivity Measurement

While an asymmetric similarity measurement such as KL-divergence may also address the issue, there is no existing work that has discussed the LSH for measuring KL-divergence. Therefore, the overhead incurred by the pairwise computation of KL-divergence is prohibitive. In this work, we propose a new measurement, called adaptivity, which can be converted to Jaccard similarity, a symmetric similarity measurement, and thus its computation can be accelerated using a novel two-level LSH (See Sec. 3.3).

The idea is to first randomly partition F_i and I into partitions of same size and then compute a set containment measurement that is defined as the ratio of the number of similar partitions shared by a training dataset and the target inference dataset to the total number of partitions in the target dataset. The formal definition of the metric is as follows:

Definition 3.1. Supposing we have two datasets $S_s = \{p_0, \dots, p_m\}$ and $S_t = \{q_0, \dots, q_n\}$, where $p_i (0 \leq i < m)$ and $q_j (0 \leq j < n)$ are partitions that have equivalent sizes (except for the last residue partition), given a threshold t , we can join S_s and S_t to identify a subset of S_t , denoted as S_t^* , which satisfies that $\forall q_j \in S_t^*, \exists p_i \in S_s$, satisfying that $D_{JS}(p_i, q_j) \leq t$. We denote the total number of partitions in S_t^* as $|S_t^*| = l$. Then we can derive a new metric to measure the adaptivity as the set containment of the target dataset (S_t) in the source dataset (S_s), denoted as: $adaptivity(S_s, S_t) = \frac{S_s \cap S_t}{S_t} \sim \frac{S_t^*}{S_t} = \frac{l}{n}$.

Based on the definition of the adaptivity metric, we further formulate the problem as illustrated in Def. 3.2, which can be easily extended to a nearest neighbor search problem of finding the top- k models that have the highest adaptivity to the target domain.

Definition 3.2. We have a database of n pre-trained models (i.e., source tasks in a source domain), represented as $M = \{M_1, \dots, M_n\}$. Each model $M_i \in M$ has a source domain S_i , which is a set of training instances that have the same types of features. We have a target query dataset, q , that is associated with a target domain T_q , which has the same types of features as source domains. Given an adaptivity threshold t' , the database should

return a set of relevant models, denoted as $M_q \subset M$, so that $\forall M_i \in M_q, adaptivity(S_i, T_q) \geq t'$.

Justification of the Adaptivity Metric. It is a common practice of leveraging partitioning to improve the accuracy of LSH in an asymmetric scenario [43]. Based on partitioning, the set containment measurement is asymmetric, indicating the ratio of the number of matched partitions to the total number of partitions in the target dataset. Thus it can effectively represent the asymmetric relationship between source domain and target domain. In addition, this set containment measurement and Jaccard similarity can be transformed to each other [43], which facilitates the determination of the threshold t' . In addition, because Minwise hash is based on Jaccard similarity (Sec. 2.2.1), we can combine the Minwise hash and the JSD-LSH (for measuring) to facilitate the computation of the adaptivity metric, as described in detail in the next section.

3.3 Novel Two-Level LSH Index

In this section, in order to accelerate the computation of the adaptivity, we propose a novel two-level LSH index that uses JSD-LSH [13] for measuring the similarity between partitions and uses Minwise LSH [8] for measuring the adaptivity from the source dataset to the target dataset.

3.3.1 JSD-LSH. We first consider the sensitiveness of JSD-LSH applied to each partition. The hash functions $\{h_{a,b}\}$, defined in Eq. 4, form a $(t, c^2 \frac{U(\lambda)}{L(\lambda)} t, e_1, e_2)$ -sensitive family for the generalized JS-divergence with parameter λ (it reduces to usual JS-divergence if $\lambda = 0.5$) [13]. This means, we can leverage $Pr[h(p_i) = h(q_j)] \geq e_1$, to determine whether $D_{JS}(p_i, q_j) \leq t$. We simply denote this mapping relationship as $e_1 = g(t)$. Please see [13] for the proof and more details. In practice, we use a family of K JSD-LSH functions by randomly choosing the hyperparameters a and b . Then we will obtain K values for each partition by applying the K functions. These K values are further partitioned into L bands. We estimate the probability that p_i and q_j are similar (i.e., $Pr[h(p_i) = h(q_j)]$) as the ratio of the matched JSD-LSH bands to L . If we see each JSD-LSH signature as a set of JSD-LSH bands, the $Pr[h(p_i) = h(q_j)]$ is actually the Jaccard similarity of the two sets. This means $jaccard(h(p_i), h(q_j)) \geq g(t)$, which connects the value of the JS-divergence threshold t as defined in Def. 3.1 and the Jaccard similarity threshold t' defined in Def. 3.2, as we explained in the next section.

3.3.2 Minwise LSH. Now we consider the adaptivity threshold t' , given source dataset $S_s = \{p_0, \dots, p_m\}$, and target dataset $S_t = \{q_0, \dots, q_n\}$, Jaccard similarity is defined as $jaccard(S_s, S_t) = \frac{S_s \cap S_t}{S_s \cup S_t} = \frac{l}{m+n-l}$, which means $jaccard(S_s, S_t) = \frac{adaptivity(S_s, S_t)}{\frac{n}{m+1} - adaptivity(S_s, S_t)}$.

Therefore, if $adaptivity(S_s, S_t) \geq t'$, it means the Jaccard similarity between source domain partitions and target domain partitions satisfies $jaccard(S_s, S_t) \geq \frac{t'}{\frac{n}{m+1} - t'}$. We leverage a family of Minwise hash functions as described in Sec. 2.2.1 to generate a Minwise hash signature for each dataset (i.e., regarded as a set of partitions), with each partition represented as a JSD-LSH signature. In addition, as mentioned at the end of Sec. 3.3.1, the equivalence of p_i and q_j is determined by $jaccard(h(p_i), h(q_j)) > g(t)$, where $h(x)$ represents the JSD-LSH function that maps x to a set of L bands, which is denoted as $h(x) = \{b_1, \dots, b_L\}$. Therefore, if $g(t) = 1/L$, we

can simplify the computation of such recursive Jaccard similarity measurement based on below equation that helps us to derive an upper-bound of the adaptivity metric:

$$jaccard(S_s, S_t) \leq L \times jaccard(\cup_{p_i} \{b_i | b_i \in p_i\}, \cup_{q_j} \{b_j | b_j \in q_j\})$$

Proof. The intuition is that if two sets of items (i.e., each item is a JSD-LSH signature computed on a partition) are similar, then if we split each item into sub-items (e.g., each sub-item is a band in a JSD-LSH signature), the two sets of sub-items are still similar to each other, and vice versa. More formally, supposing $\exists A \subset S_s$ satisfying $\forall p_i \in A, \exists q_j \in S_t$, and $jaccard(h(p_i), h(q_j)) \geq 1/L$, we have $jaccard(S_s, S_t) = \frac{|A|}{|S_s \cup S_t|}$. Additionally, $jaccard(h(p_i), h(q_j)) \geq 1/L$ indicates that $\exists b_i \in p_i$ and $b_j \in q_j$, satisfying $b_i = b_j$. Therefore, $jaccard(\cup_{p_i} \{b_i | b_i \in p_i\}, \cup_{q_j} \{b_j | b_j \in q_j\}) \geq \frac{|A|}{|S_s \cup S_t| \times L}$. Therefore, $jaccard(S_s, S_t) \leq L \times jaccard(\cup_{p_i} \{b_i | b_i \in p_i\}, \cup_{q_j} \{b_j | b_j \in q_j\})$, and we can use the latter as an upper-bound estimation of the former metric.

Based on the approximation, a JSD-LSH signature is first generated for each partition and each signature is further split into L bands. Then the set of partitions in a dataset will be flattened into a set of JSD-LSH bands. Given user-specified adaptivity threshold t' , we can decide whether two sets have their Jaccard similarity higher than $\frac{t'}{\frac{n}{m}+1-t'}$ times a factor of $L = \frac{1}{g(t)}$ by comparing their Minwise hash signatures as described in Sec. 2.2.1.

One potential problem is that the value of n is probably not a constant shared by all training datasets. To address the problem, we leverage a padding strategy as in asymmetric Minwise LSH [33]. We select an upper bound of n , denoted as n_u , then we pad every training dataset to have n_u partitions by appending $(n_u - n)$ JSD-LSH signatures for non-existing partitions. The values in these padding signatures will not exist in real signatures, so the padding will not affect the precision, while it may lead to false negatives.

The algorithm based on the two-level LSH index to address the problem formulated in Def. 3.2 is illustrated in Alg. 1.

3.4 User Interfaces Design

Our system allows users to upload models by providing the path to the model files, the path to its training dataset, as well as the model metadata. Note that the training dataset will not be directly stored in the system. Instead, the training dataset will be partitioned and scanned to generate the JSD-LSH signatures and further the Minwise LSH signatures, and only these signatures will be stored in the system. Then the Minwise LSH signatures will be split into many bands, and each band is inserted into a corresponding hash table, which is transparent to the users. An example of using the system is illustrated in Listing. 1.

Listing 1: Illustration of User Interfaces

```
class ModelHandler:
    #return predictions
    def execute(input)
class AppBase:
    #return a set of ModelHandler instances based on the target data
    def search(path_to_target_data, general=false)
    #load models during setup time, implemented by user
```

Algorithm 1 Model discovery algorithm

```
1: INPUT1:  $\mathcal{T}_q$  (query dataset, i.e., the target dataset)
2: INPUT2:  $t'$  (the adaptivity threshold)
3: INPUT3:  $L$  (the number of bands in each JSD-LSH hash)
4: INPUT4:  $L_m$  (the number of bands in each Minwise hash)
5: INPUT5:  $\{hashmap_i\} (1 \leq i \leq L_m)$  (one hashmap for each band of the Minwise-LSH.)
6: OUTPUT:  $\mathcal{M}_q$ 
7:  $m \leftarrow |\mathcal{T}_q|$ 
8:  $t^* \leftarrow L \times \frac{t'}{\frac{n}{m}+1-t'}$ 
9:  $\mathcal{Q} \leftarrow partition(\mathcal{T}_q)$ 
10:  $S \leftarrow \phi$ 
11: for  $q \in \mathcal{Q}$  do
12:    $S \leftarrow S \cup JSDLSH(q)$ 
13: end for
14:  $S' \leftarrow MinwiseLSH(S)$  ( $S'$  consists of  $L_m$  bands)
15:  $\mathcal{M}, \mathcal{M}_q \leftarrow \phi$ 
16: for  $s_i \in S'$  do
17:    $\mathcal{M}_i \leftarrow hashmap_i.lookup(s_i)$ 
18:    $\mathcal{M} \leftarrow \mathcal{M} \cup \mathcal{M}_i$ 
19: end for
20: for  $M_i \in \mathcal{M}$  do
21:   if  $\frac{M.count(M_i)}{L_m} > t^*$  then
22:      $\mathcal{M}_q \leftarrow \{m\} \cup \mathcal{M}_q$ 
23:   end if
24: end for
25: return  $\mathcal{M}_q$ 
```

```
def setup()
    #process the target data, implemented by user
    def process(path_to_target_data)
class MyApp(AppBase):
    def setup(self):
        self.models = self.search("./data/train/myData")
    def process(self, "./data/train/myData"):
        for (m in self.models):
            results.append(m.execute(myData))
        return results
```

4 EMPIRICAL EVALUATION

In this section, we focus on answering following questions:

- (1) How effective is the proposed adaptivity metric, compared to other similarity metrics? (Tab. 3, Tab. 4, Tab. 6, Tab. 7, Tab. 9, Tab. 10, Tab. 11, Tab. 12)
- (2) How will the two-level LSH index accelerate the model discovery process, and how will it affect the accuracy, compared to other adaptivity computation strategies? (Tab. 5, Tab. 8)
- (3) How effective is using JSD-LSH to estimate the JS-divergence? (Fig. 2)

4.1 Workloads and Datasets

We evaluate our proposed methodology using four workloads.

1. Activity Recognition. Human activity recognition (HAR), is to predict the activities (e.g., walking, sitting, running, lying) based on data collected from multiple sensors attached to human body. HAR is a hot research topic in the pervasive computing area, and has been widely applied to indoor localization, sleep state detection, smart home sensing, and virtual reality [6]. In this work, we use three

public activity recognition datasets, including OPPORTUNITY [11], PAMAP2 [31], and DSADS [7]. The OPPORTUNITY dataset is collected from four human subjects executing various activities with sensors attached to more than five body parts. The PAMAP2 dataset is collected from nine subjects performing 18 activities with sensors attached to three body parts. The DSADS dataset is collected from eight subjects wearing sensors on five body parts. A dataset is collected for each body part, therefore there are 13 datasets in total. Each dataset has 81 features [36] and 1920 to 5022 samples. Then we apply our approach to a series of scenarios where models are trained on each of these tables, and given a target dataset collected from a subject on a specific body part, we want to select a model to achieve the best accuracy on the target dataset.

2. Image Recognition. We randomly sample images from the Cifar-10 dataset without replacement to create five image datasets with distinct probability distributions. Cifar-10 is a collection of images with labels of ten classes, with each containing 5,000 images. Then we make four of the subsets skewed by adding 5,000 images of the fourth class to the second and the fourth subset; 5,000 images of the second class to the third subset; 5,000 images of the eighth class to the third subset. These five datasets are called as Balanced, Skewed-1, Skewed-2, Skewed-3, Skewed-4 correspondingly. To evaluate the effectiveness of our proposed metrics for the image recognition scenario, we train ResNet56v1 model on each of the subsets respectively, and then perform five experiments of searching for the best model to be reused on each dataset.

3. Entity Matching (EM). This task is to decide *whether two tuples are referring to the same entity* [9, 16, 17, 24, 41]. We mainly apply DeepMatcher [29], which is an EM tool based on deep learning, to four datasets¹: Walmart-Amazon, Abt-Buy, DBLP-Scholar, DBLP-ACM; and 11 smaller datasets from the Magellan Data Repository [14]². Each task contains training and testing samples collected from two different datasets. For example, IMDB-TMD is to match movie tuples collected from IMDB and TMD respectively, and IMDB-RottenTomatoes is to match movie tuples from IMDB and Rotten Tomatoes respectively. We focus on the model selection scenarios where a dataset is used for query, and a set of models are trained for the EM tasks on the rest of the datasets as candidates, and then we try to select top- k candidate models to serve the query dataset.

4. Natural Language Processing (NLP). We are mainly interested in two types of NLP tasks. The first task is to identify whether sentence pairs have equivalent semantic meanings. For this task, we train models on three different datasets: Microsoft Research Paraphrase Corpus (MRPC) that includes 3,549 samples; Quora Question Pairs (QQP) that consists of 363,192 samples; and Paraphrase Adversaries from Word Scrambling (PAWS) which has 49,401 samples. The second task is about natural language inference (NLI), which is to identify the textual entailment relationship [26] between sentence pairs. Similar to the first task, we train models on three different datasets: Recognizing Textual Entailment (RTE) that has 2,490 samples; Question NLI, which contains 510,711 samples; and the SCITAIL dataset including 5,302 samples, which is an entailment dataset created from multiple-choice science exams and

web sentences. We train models for these tasks using the same pre-trained BERT base model³, denoted as uncased_L-12_H-768_A-12, which contains 12 layers and 110 millions of parameters. Then for each task, we use one of the three datasets as the target dataset, and try to choose the best model to serve on the target dataset.

4.2 Evaluation Methodology

One major evaluation task is to compare our proposed adaptivity metric to other alternative similarity measurements, including: (1) JS-divergence (a special case of adaptivity when each dataset has only one partition), which is a similarity measurement for two probability distributions, as mentioned earlier in the paper. (2) L2-distance, which is to get the center of the instances by averaging all feature vectors for the source and target datasets, and then compute euclidean distance between two centers; (3) Source accuracy, which is the accuracy of the candidate model on its training dataset. We consider this metric, as this information is widely available in most types of model databases and in practice, users usually rely on this information for model discovery.

To compare the effectiveness of these metrics, we use:

- (1) **Pearson correlation coefficient (PCC)** [23], which is a measure of the linear correlation coefficient between two random variables, to show the correlation of various metrics to the target accuracy (i.e., the prediction accuracy of the candidate model on the target dataset).
- (2) **Top- k -error**, defined as the number of wrong predictions that fail to correctly select the optimal model in the top- k results (ranking of the k selected models is not important here) to the total number of predictions. For example, for top-3 search, if the ground truth is {C}, the prediction of {B, C, D} is a correct prediction. We also use the error rate metric for evaluating the effectiveness of different strategies of computing the adaptivity metric. In addition, for the comparison of different strategies of computing the adaptivity, we also evaluate the query latency for a model discovery process. We tune hyperparameters including the value of r for JSD-LSH, the number of hash functions and the number of bands for Minwise Hash and JSD-LSH, and the number of partitions for each test case by searching the combination of hyperparameters in a pre-specified range.

4.3 Results

4.3.1 Activity Recognition. We first evaluate the Activity Recognition workload for the effectiveness of the proposed adaptivity metric, JSD-LSH, and the two-level LSH index.

Effectiveness of the Adaptivity Metric. The comparison of Pearson correlation coefficient values of various metrics, including the adaptivity, the JS-divergence, the L2-distance, and the source accuracy, to the target accuracy are illustrated in Tab. 3. It shows that our proposed adaptivity metric achieves a better correlation with target accuracy compared to other metrics. Also among the four metrics, the source accuracy exhibits the least relevance with the target accuracy, which verifies our assumption that the similarity between the source and the target plays an important role in determining the target accuracy (i.e., model adaptivity to the target domain). We also compute the top-1, top-2, top-3 error rates, as

¹<https://github.com/anhaidgroup/deepmatcher/blob/master/Datasets.md>

²<https://sites.google.com/site/anhaidgroup/useful-stuff/data>

³https://tfhub.dev/google/bert_uncased_L-12_H-768_A-12/1

illustrated in Tab. 4. As mentioned, Top- k error rate is defined to be the ratio of the number of searches that failed to return all k models correctly to the total number of searches. The results show that the adaptivity works better than other metrics.

Table 3: Comparison of Pearson correlation coefficient (PCC) for activity recognition (the best Pearson correlation coefficient values are highlighted in bold).

target	pcc of adaptivity	pcc of JS-divergence	pcc of l2-distance	pcc of source-accuracy
dsads_la	0.61	0.17	0.54	0.44
dsads_ll	0.72	0.40	0.68	0.40
dsads_ra	0.51	0.13	0.54	0.37
dsads_rl	0.58	0.49	0.75	0.58
dsads_t	0.19	0.11	0.42	0.44
oppo_b	0.89	0.66	0.77	0.45
oppo_ll	0.79	0.54	0.54	0.34
oppo_lua	0.91	0.72	0.76	0.48
oppo_rla	0.76	0.54	0.52	0.28
oppo_rua	0.84	0.67	0.70	0.44
pamap_a	0.69	0.68	0.54	0.24
pamap_c	0.77	0.69	0.70	0.06
pamap_w	0.83	0.74	0.63	0.14

Effectiveness of the Two-Level Index. Now we compare the accuracy and efficiency of using different adaptivity computation strategies including pair-wise comparison, two-level LSH without flattening, and two-level LSH index with flattening. Flattening refers to the step that flattens a set of partitions into a set of JSD-LSH bands, which is the approach that we leveraged for approximating the recursive Jaccard similarity as justified in Sec. 3.3.2. The case where flattening is not applied is equivalent to the case where flattening is applied but the number of JSD-LSH bands L is set to 1. The results are illustrated in Tab. 5. Here, for JSD-LSH, we use $r = 1.4$, number of bands is $L = 200$, number of hash functions is $K = 800$. For Minwise hash, we use 256 minwise hash functions, and use 128 bands. We also allocate 55 samples in each partition, we can see that if without the flattening step, the accuracy will significantly drop. In addition, the two-level index achieves about 65 \times speedup compared to the pair-wise approach for this scenario. **Effectiveness for JSD-LSH.** To justify the use of JSD-LSH, which is a relatively new LSH function and the first LSH function for indexing JS-divergence with a bound, we compare the accuracy and latency of LSH for JS-divergence to a baseline approach that computes the JS-divergences between the target dataset and each of the source datasets directly without using LSH techniques. For this experiment, we created 162 tables by sampling these 13 activity recognition datasets, so that each table represents samples collected from a body part for specific subject. The results are illustrated in Fig. 2, in which each bar represents a test case that uses one

Table 4: Comparison of different similarity measurements for activity recognition

	adaptivity	JS-divergence	l2-distance	source-accuracy
top-1 error	15%	31%	23%	92%
top-2 error	0%	0%	0%	92%
top-3 error	0%	0%	0%	77%

Table 5: Comparison of adaptivity computation strategies for activity recognition

	top-1 error	top-2 error	top-3 error	latency (milli-sec)
pair-wise	15%	0%	0%	24,627
2-level LSH w/o flattening	54%	62%	54%	377
2-level LSH w/ flattening	31%	8%	0%	377

Table 6: Comparison of Pearson correlation coefficient for image recognition (the best Pearson correlation coefficient values are highlighted in bold)

target	pcc of adaptivity	pcc of JS-divergence	pcc of l2-distance	pcc of source-accuracy
Balanced	0.93	0.92	0.62	0.80
Skewed-1	0.86	0.53	0.52	0.78
Skewed-2	0.39	0.70	0.66	0.07
Skewed-3	0.25	0.06	0.18	0.73
Skewed-4	0.77	0.41	0.38	0.39

table as the query (i.e., target dataset), and the rest of the tables as the sources. In this experiment, for each test, we assume the source datasets that have the JS-divergence with the query dataset smaller than 0.1 compose the ground truth. We observe that using LSH can significantly accelerate the overall latency required for JS-divergence comparison, and achieve 5 \times speedup on average, while the precision of our proposed approach is 100% and recall is above 93%.




Figure 2: Comparison of LSH for JS-divergence (denoted as LSH) to a pair-wise computation of JS-divergence (denoted as baseline) for 162 tables sampled from 13 activity recognition datasets.

4.3.2 Image Recognition. In this scenario, for each Cifar-10 image, we normalize all pixel values by dividing each value by 255, and then flatten the image from a tensor of the shape (32, 32, 3) into a one-dimensional vector of the shape (1, 3072). Then, we compute the adaptivity, JS-divergence and L2-distance on the flattened dataset.

Effectiveness of the Adaptivity Metric. We evaluate and compare the effectiveness of adaptivity, JS-divergence, L2-distance, and source accuracy for searching the best model to serve on a target dataset. The Pearson correlation coefficient of each metric to the target accuracy in each experiment is illustrated in Tab. 6. The top-1 and top-2 error rate for searching the best model for each experiment is illustrated in Tab. 7. We only show top-1 and top-2 error rate, because we have only four candidate models in this case. **Effectiveness of the Two-Level Index.** We also compare the effectiveness of using the two-level index to the pair-wise strategy for computing the adaptivity metric, as illustrated in Tab. 8. For this scenario each partition has 200 samples. The two-level index

Table 7: Comparison of error rate for image recognition

	adaptivity	JS-divergence	l2-distance	source-accuracy
top-1 error	20%	40%	40%	60%
top-2 error	20%	0%	0%	40%

Table 8: Comparison of adaptivity computation strategies for image recognition

	top-1 error	top-2 error	latency (milli-sec)
pair-wise	20%	0%	14,585
2-level LSH w/ flattening	0%	0%	1,652

Table 9: Comparison of Pearson correlation coefficient of various metrics to the target accuracy for entity matching (the best Pearson correlation coefficient values are highlighted in bold)

target	pcc of adaptivity	pcc of JS-divergence	pcc of source-accuracy
Abt_Buy	0.99	0.99	0.99
Dplp_Acm	0.92	0.94	0.70
Dbip_Scholar	0.91	0.94	0.83
Walmart_Amazon	0.99	0.99	0.73
MyAnimeList_AnimePlanet	0.61	0.53	0.51
Bikedekho_Bikewale	0.51	0.16	0.19
Amazon_Barnes	0.24	0.17	0.22
GoodReads_Barnes	0.61	0.53	0.20
Barnes_Half	0.29	0.44	0.31
RottenTomatoes_IMDB	0.11	0.34	0.16
IMDB_TMD	0.49	0.67	0.44
IMDB_RottenTomatoes	0.51	0.75	0.22
Amazon_RottenTomatoes	0.60	0.64	0.38
RogerElbert_IMDB	0.58	0.13	0.07
YellowPages_Yelp	0.82	0.67	0.44

Table 10: Comparison of error rate for entity matching

	adaptivity	JS-divergence	source-accuracy
top-1 error	13%	47%	33%
top-2 error	0%	27%	27%

is observed to achieve more than 8× speedup, while achieving similar accuracy compared to the pair-wise strategy. The speedup is relatively smaller than the activity recognition scenario. That's mainly because we involve fewer number of datasets and fewer number of partitions in this scenario.

4.3.3 Entity Matching. In this experiment, different from activity recognition, most attributes contain text-based values. We represent the training dataset of each entity matching task as a bag of words over a shared dictionary for computing the JS-divergence and adaptivity. We compare the Pearson correlation coefficient values of the JS-divergence metric, the adaptivity metric, and source accuracy, to the target accuracy for each EM task, as illustrated in Tab. 9. We also compare the overall accuracy in terms of top1-error and top2-error for all fifteen tasks, as illustrated in Tab. 10. The results show that the adaptivity metric outperforms other metrics in selecting the models to serve with the best accuracy.

4.3.4 Natural Language Processing. For this experiment, as mentioned in Sec. 4.1, because each task involves merely three datasets, a model discovery scenario for a target dataset only requires to compare two datasets. Therefore, instead of computing the Pearson coefficient for each search scenario, we choose to compute the Pearson coefficient between the metrics and the overall accuracy of each of two tasks. Each task consists of three different model discovery scenario. For the same reason, we only consider top-1 error for this experiment, which is counted over all six model discovery scenarios across the two tasks. The results are illustrated

Table 11: Comparison of Pearson correlation coefficient for NLP (the best Pearson correlation coefficient values are highlighted in bold)

target	pcc of adaptivity	pcc of JS-divergence	pcc of source-accuracy
Task1	0.61	0.71	0.02
Task2	0.76	0.87	0.10

Table 12: Comparison of error rate for NLP

	adaptivity	JS-divergence	source-accuracy
top-1 error	0%	16.7%	33.3%

in Tab. 11 and Tab. 12, which show that while adaptivity has less correlation with the target accuracy compared to JS-divergence, it can more effectively predict the best model for serving with zero error rate for all six search scenarios. The source accuracy performs significantly worse than other metrics.

5 RELATED WORKS

In recent, numerous works are proposed to address open data discovery problems, including automatically discover table unionability [30] and joinability [42, 43], and related tables [40]. Most of these works are trying to identify all similar features using LSH techniques based on Jaccard similarity or variants. While these works can be leveraged to accelerate the match of JSD-LSH signatures, they are not directly applicable in selecting related models, because the existence of a significant portion of shared features between the source and target datasets doesn't mean the probability distributions over the two feature spaces are similar.

Zamir and et al. [39] propose a computational approach to find the transferable relationships, abstracted as taxonomy among the different computation vision tasks, e.g., depth estimation, edge detection, point matching and etc., so that some tasks can be trained using other tasks' output and thus require less training data and supervision budget. Their work is focused on the adaptivity from a task's output domain to another task's input domain. In contrast, our work is focused on the adaptivity of tasks' input domains. Also while their work is limited to computer vision, our work is targeting a more generalized scenario. Kornblith and et al. [20] study the relationship between the ImageNet model architectures such as ResNet, MobileNet, Inception-ResNet, etc., and the transfer learning performance. They observe that transfer learning based on fine-tune techniques (i.e., just tune the last two layers) even using a "better" model architecture may not achieve "better" accuracy. This means in regard of the effectiveness of transfer learning, the divergence of features between the source and target domains is more important than the difference in model architectures. Their observation indicates that our work may be applicable to a more general scenario for finding related models that have inconsistent architectures for reuse.

6 CONCLUSIONS

In a broad class of AI applications, such as Uber's marketplace-level forecasting and downtime prediction for manufacturers, a large number of model versions will be created for different locations and different time points. To facilitate model reuse, in this work, we systematically explore the problem of searching models from multiple versions of models that have similar architectures and

training methodologies, for serving on a target query dataset. To address the problem, we propose a novel two-level LSH index based on a new adaptivity metric that is not only asymmetric but also can be converted into Jaccard similarity. We demonstrate the accuracy and efficiency of the proposed system through extensive experiments on a number of workloads including activity recognition, entity matching, image recognition, and natural language processing. Our proposed work can greatly save the human costs in model searching and shorten the deep learning model deployment time for production and research.

REFERENCES

- [1] [n. d.]. *Deep Learning Hub*. <https://dlhub.app/>.
- [2] [n. d.]. *Impact of slow page load time on website performance*. <https://medium.com/@vikiigreen/impact-of-slow-page-load-time-on-website-performance-40d5c9ce568a>.
- [3] [n. d.]. *PyTorch Hub*. <https://pytorch.org/hub/>.
- [4] [n. d.]. *Tensorflow Hub*. <https://www.tensorflow.org/hub>.
- [5] [n. d.]. *Uber Open Summit 2018*. <https://uberopen2018.splashthat.com/>.
- [6] Akin Avci, Stephan Bosch, Mihai Marin-Perianu, Raluca Marin-Perianu, and Paul Havinga. 2010. Activity recognition using inertial sensing for healthcare, wellbeing and sports applications: A survey. In *23th International conference on architecture of computing systems 2010*. VDE, 1–10.
- [7] Billur Barshan and Murat Cihan Yüksek. 2014. Recognizing daily and sports activities in two open source machine learning environments using body-worn sensor units. *Comput. J.* 57, 11 (2014), 1649–1667.
- [8] Andrei Z Broder. 1997. On the resemblance and containment of documents. In *Proceedings. Compression and Complexity of SEQUENCES 1997 (Cat. No. 97TB100171)*. IEEE, 21–29.
- [9] Riccardo Cappuzzo, Paolo Papotti, and Saravanan Thirumuruganathan. 2020. Creating embeddings of heterogeneous relational datasets for data integration tasks. In *Proceedings of the 2020 ACM SIGMOD International Conference on Management of Data*. 1335–1349.
- [10] Moses S Charikar. 2002. Similarity estimation techniques from rounding algorithms. In *Proceedings of the thirty-fourth annual ACM symposium on Theory of computing*. 380–388.
- [11] Ricardo Chavarriaga, Hesam Sagha, Alberto Calatroni, Sundara Tejaswi Digu-marti, Gerhard Tröster, José del R Millán, and Daniel Roggen. 2013. The Opportunity challenge: A benchmark database for on-body sensor-based activity recognition. *Pattern Recognition Letters* 34, 15 (2013), 2033–2042.
- [12] Andrew Chen, Andy Chow, Aaron Davidson, Arjun DCunha, Ali Ghodsi, Sue Ann Hong, Andy Konwinski, Clemens Mewald, Siddharth Murching, Tomas Nykodym, et al. 2020. Developments in MLflow: A System to Accelerate the Machine Learning Lifecycle. In *Proceedings of the fourth international workshop on data management for end-to-end machine learning*. 1–4.
- [13] Lin Chen, Hossein Esfandiari, Gang Fu, and Vahab Mirrokni. 2019. Locality-Sensitive Hashing for f-Divergences: Mutual Information Loss and Beyond. In *Advances in Neural Information Processing Systems*. 10044–10054.
- [14] Sanjib Das, AnHai Doan, Paul Suganthan G. C., Chaitanya Gokhale, Pradap Konda, Yash Govind, and Derek Paulsen. [n. d.]. The Magellan Data Repository. <https://sites.google.com/site/anhaidgroup/projects/data>.
- [15] Mayur Datar, Nicole Immorlica, Piotr Indyk, and Vahab S Mirrokni. 2004. Locality-sensitive hashing scheme based on p-stable distributions. In *Proceedings of the twentieth annual symposium on Computational geometry*. 253–262.
- [16] Muhammad Ebraheem, Saravanan Thirumuruganathan, Shafiq Joty, Mourad Ouzzani, and Nan Tang. 2017. DeepER–Deep Entity Resolution. *arXiv preprint arXiv:1710.00597* (2017).
- [17] Raul Castro Fernandez, Essam Mansour, Abdulhakim A Qahtan, Ahmed Elmagarmid, Ihab Ilyas, Samuel Madden, Mourad Ouzzani, Michael Stonebraker, and Nan Tang. 2018. Seeping semantics: Linking datasets using word embeddings for data discovery. In *2018 IEEE 34th International Conference on Data Engineering (ICDE)*. IEEE, 989–1000.
- [18] Carlos Roberto González Palacios. 2015. *Optimal Data Distributions in Machine Learning*. Ph.D. Dissertation. California Institute of Technology.
- [19] Piotr Indyk and Rameez Motwani. 1998. Approximate nearest neighbors: towards removing the curse of dimensionality. In *Proceedings of the thirtieth annual ACM symposium on Theory of computing*. 604–613.
- [20] Simon Kornblith, Jonathon Shlens, and Quoc V Le. 2019. Do better imagenet models transfer better?. In *Proceedings of the IEEE conference on computer vision and pattern recognition*. 2661–2671.
- [21] Solomon Kullback and Richard A Leibler. 1951. On information and sufficiency. *The annals of mathematical statistics* 22, 1 (1951), 79–86.
- [22] Lillian Lee. 2000. Measures of distributional similarity. *arXiv preprint cs/0001012* (2000).
- [23] Joseph Lee Rodgers and W Alan Nicewander. 1988. Thirteen ways to look at the correlation coefficient. *The American Statistician* 42, 1 (1988), 59–66.
- [24] Yuliang Li, Jinfeng Li, Yoshihiko Suhara, AnHai Doan, and Wang-Chiew Tan. 2020. Deep Entity Matching with Pre-Trained Language Models. *arXiv preprint arXiv:2004.00584* (2020).
- [25] Jianhua Lin. 1991. Divergence measures based on the Shannon entropy. *IEEE Transactions on Information theory* 37, 1 (1991), 145–151.
- [26] Bill MacCartney and Christopher D Manning. 2009. *Natural language inference*. Citeseer.
- [27] Xian-Ling Mao, Bo-Si Feng, Yi-Jing Hao, Liqiang Nie, Heyan Huang, and Guihua Wen. 2017. S2JSD-LSH: A locality-sensitive hashing schema for probability distributions. In *Proceedings of the Thirty-First AAAI Conference on Artificial Intelligence*. 3244–3251.
- [28] Hui Miao, Ang Li, Larry S Davis, and Amol Deshpande. 2017. Modelhub: Deep learning lifecycle management. In *2017 IEEE 33rd International Conference on Data Engineering (ICDE)*. IEEE, 1393–1394.
- [29] Sidharth Mudgal, Han Li, Theodoros Rekatsinas, AnHai Doan, Youngchoon Park, Ganesh Krishnan, Rohit Deep, Esteban Arcaute, and Vijay Raghavendra. 2018. Deep learning for entity matching: A design space exploration. In *Proceedings of the 2018 International Conference on Management of Data*. 19–34.
- [30] Fatemeh Nargesian, Erkang Zhu, Ken Q Pu, and Renée J Miller. 2018. Table union search on open data. *Proceedings of the VLDB Endowment* 11, 7 (2018), 813–825.
- [31] Attila Reiss and Didier Stricker. 2012. Introducing a new benchmarked dataset for activity monitoring. In *2012 16th International Symposium on Wearable Computers*. IEEE, 108–109.
- [32] Sebastian Schelter, Joos-Hendrik Boese, Johannes Kirschnick, Thoralf Klein, and Stephan Seufert. 2017. Automatically tracking metadata and provenance of machine learning experiments. In *Machine Learning Systems Workshop at NIPS*. 27–29.
- [33] Anshumali Shrivastava and Ping Li. 2015. Asymmetric minwise hashing for indexing binary inner products and set containment. In *Proceedings of the 24th international conference on world wide web*. 981–991.
- [34] Chong Sun, Nader Azari, and Chintan Turakhia. 2020. Gallery: A Machine Learning Model Management System at Uber. In *EDBT*. 474–485.
- [35] Manasi Vartak, Harihar Subramanyam, Wei-En Lee, Srinidhi Viswanathan, Saadiyah Husnoo, Samuel Madden, and Matei Zaharia. 2016. ModelDB: a system for machine learning model management. In *Proceedings of the Workshop on Human-In-the-Loop Data Analytics*. 1–3.
- [36] Jindong Wang, Yiqiang Chen, Lisha Hu, Xiaohui Peng, and S Yu Philip. 2018. Stratified transfer learning for cross-domain activity recognition. In *2018 IEEE International Conference on Pervasive Computing and Communications (PerCom)*. IEEE, 1–10.
- [37] Christian Weber, Pascal Hirmer, and Peter Reimann. 2020. A Model Management Platform for Industry 4.0–Enabling Management of Machine Learning Models in Manufacturing Environments. In *International Conference on Business Information Systems*. Springer, 403–417.
- [38] Yang Yang, De-Chuan Zhan, Ying Fan, Yuan Jiang, and Zhi-Hua Zhou. 2017. Deep Learning for Fixed Model Reuse. In *AAAI*. 2831–2837.
- [39] Amir R Zamir, Alexander Sax, William Shen, Leonidas J Guibas, Jitendra Malik, and Silvio Savarese. 2018. Taskonomy: Disentangling task transfer learning. In *Proceedings of the IEEE conference on computer vision and pattern recognition*. 3712–3722.
- [40] Yi Zhang and Zachary G Ives. 2020. Finding Related Tables in Data Lakes for Interactive Data Science. In *Proceedings of the 2020 ACM SIGMOD International Conference on Management of Data*. 1951–1966.
- [41] Chen Zhao and Yeye He. 2019. Auto-EM: End-to-end Fuzzy Entity-Matching using Pre-trained Deep Models and Transfer Learning. In *The World Wide Web Conference*. 2413–2424.
- [42] Erkang Zhu, Dong Deng, Fatemeh Nargesian, and Renée J Miller. 2019. Josie: Overlap set similarity search for finding joinable tables in data lakes. In *Proceedings of the 2019 International Conference on Management of Data*. 847–864.
- [43] Erkang Zhu, Fatemeh Nargesian, Ken Q Pu, and Renée J Miller. 2016. LSH ensemble: Internet-scale domain search. *arXiv preprint arXiv:1603.07410* (2016).